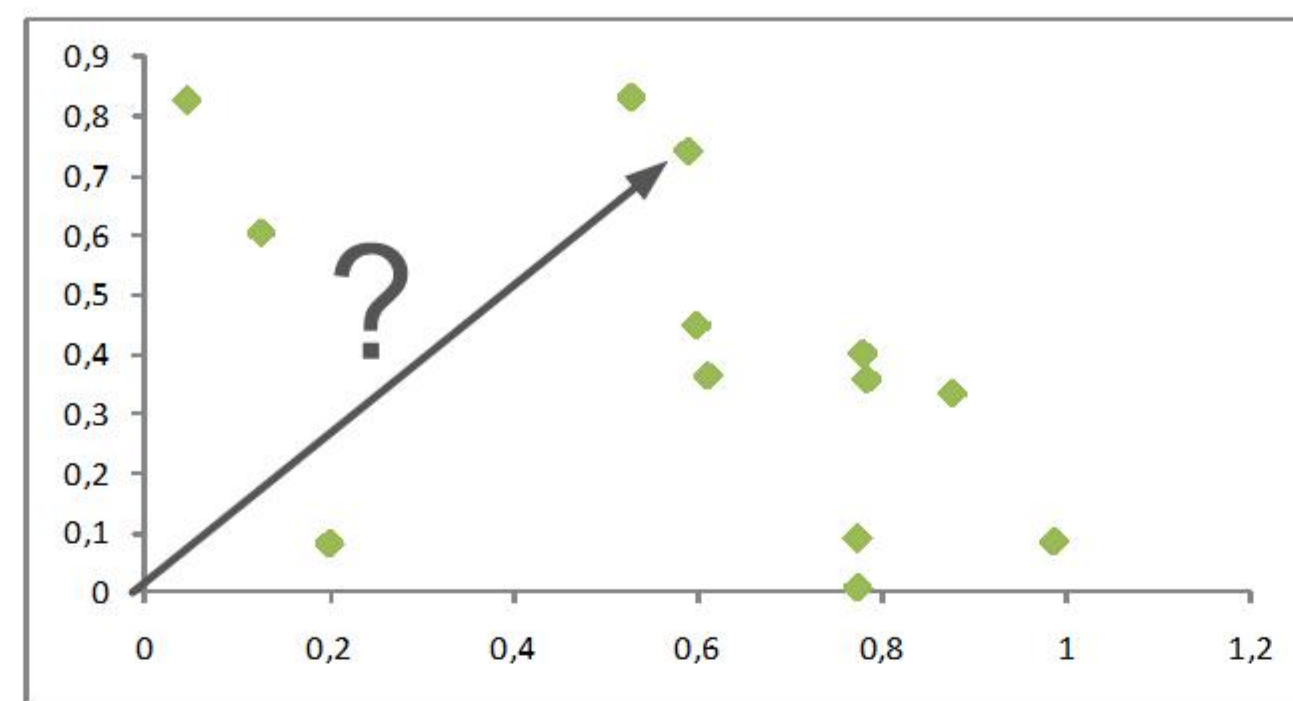




Goal: Use massively parallel hardware for core database operations


## Multidimensional Search

```
SELECT *
FROM table
WHERE x = 0.6
AND y BETWEEN 0.6 AND 0.8
```




### Methods

Room filling curves



Concatination of dimension

Multidimensional index



Sequential algorithms and data structures


**Current principle:**  
Save CPU Cycles

Speculative Execution as more efficient approach for GPU based data management?




Idle Cores  
are the  
Devil's Tools  
Don't wait for Work!

## Speculative Multidimensional Index



### Evaluation



Execution on a c2050 vs. Intel Nehalem 3,5GHz

## Conclusion and Outlook


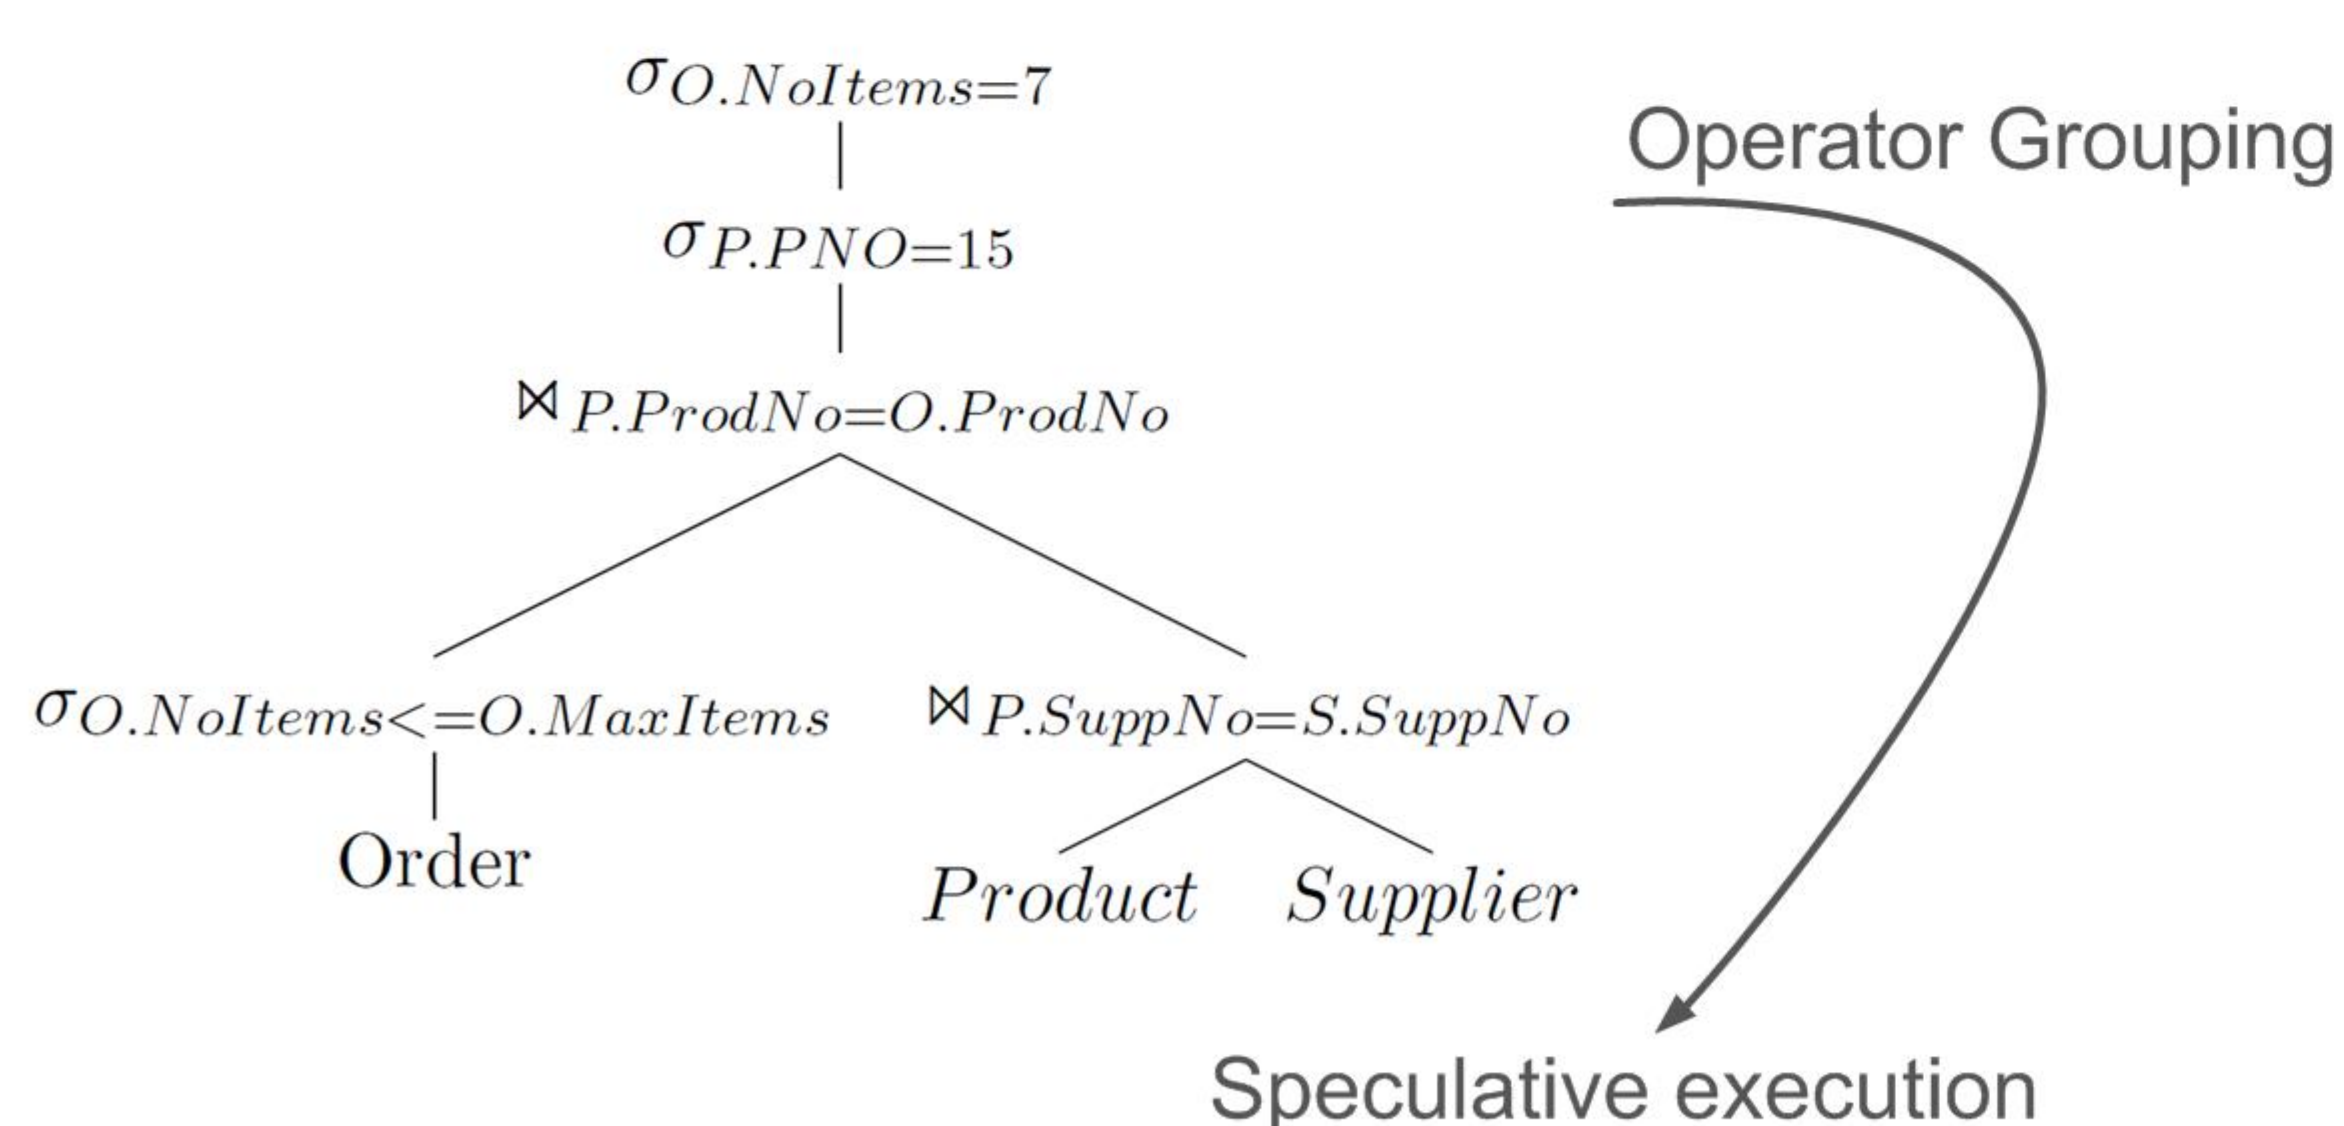
### Speculative data management....

- ....works on a fine grain level
- ....helps to search for data faster
- ....requires massively parallel hardware such as GPUs

### Leverage speculative data management....

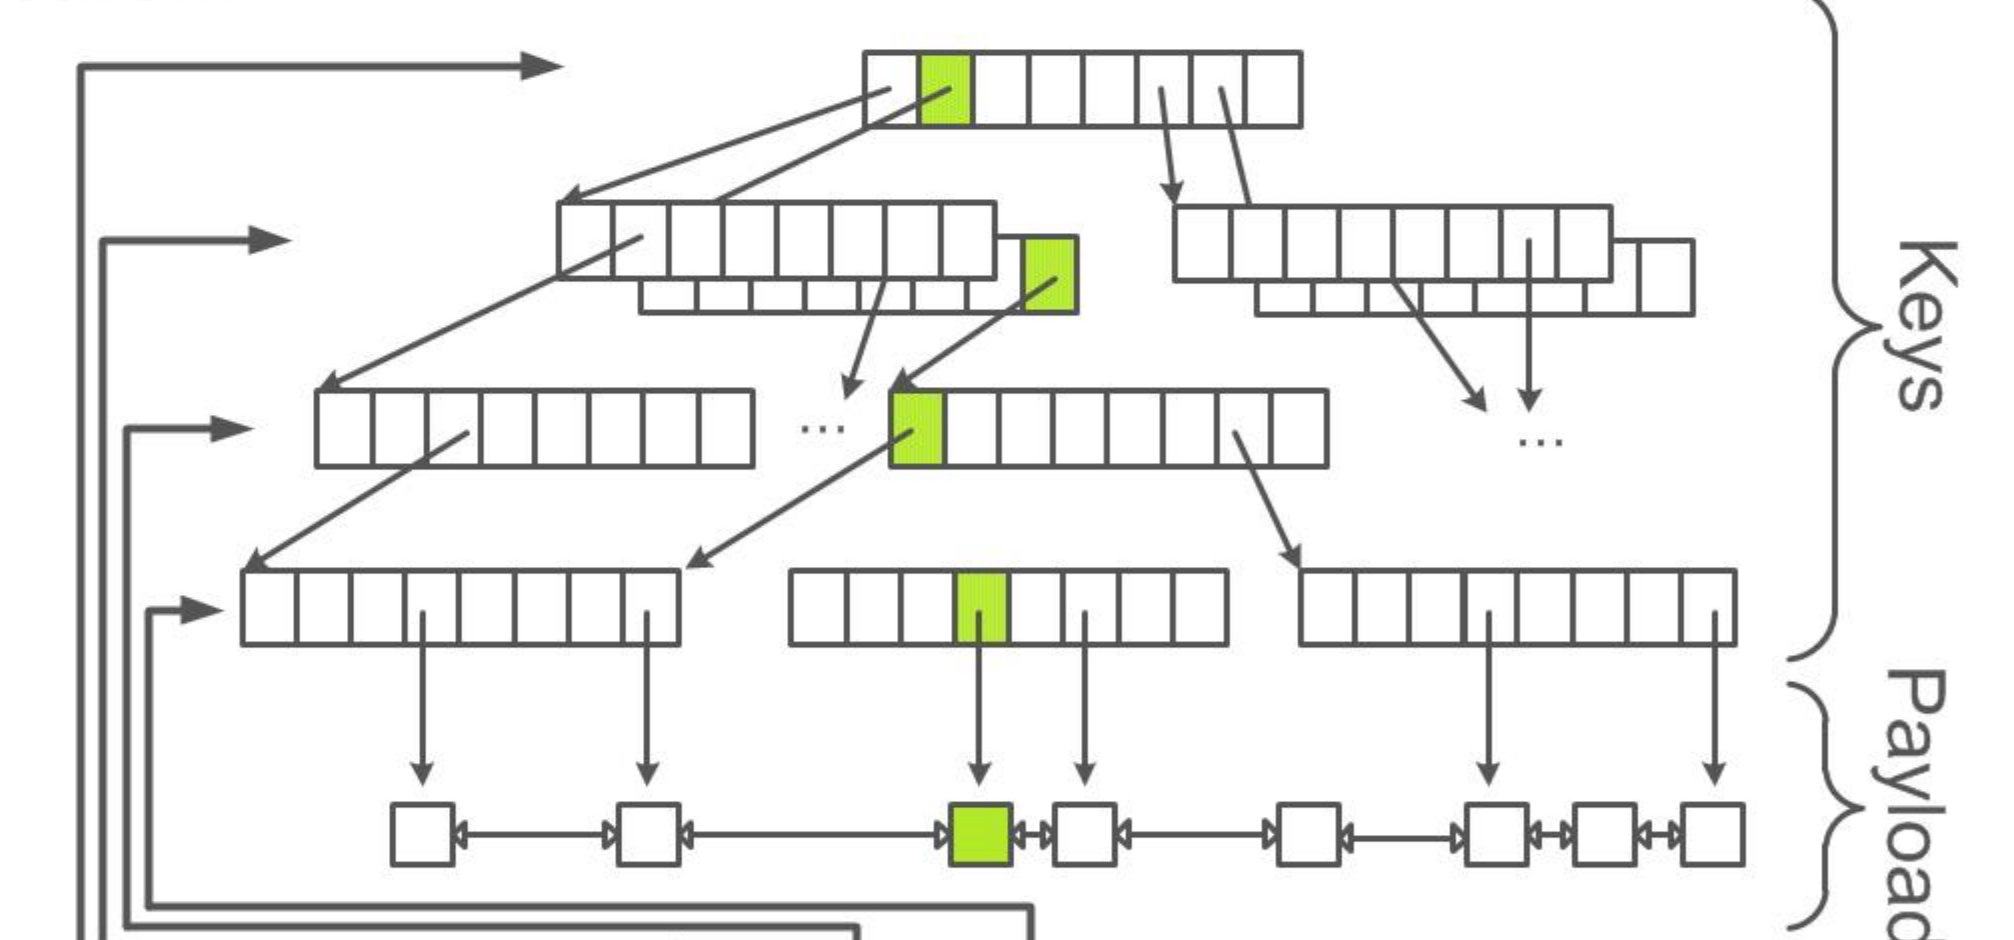
- ....on a higher level (e.g. execution plans)
- ....with multiple massively parallel devices

Parallel Query Plan:

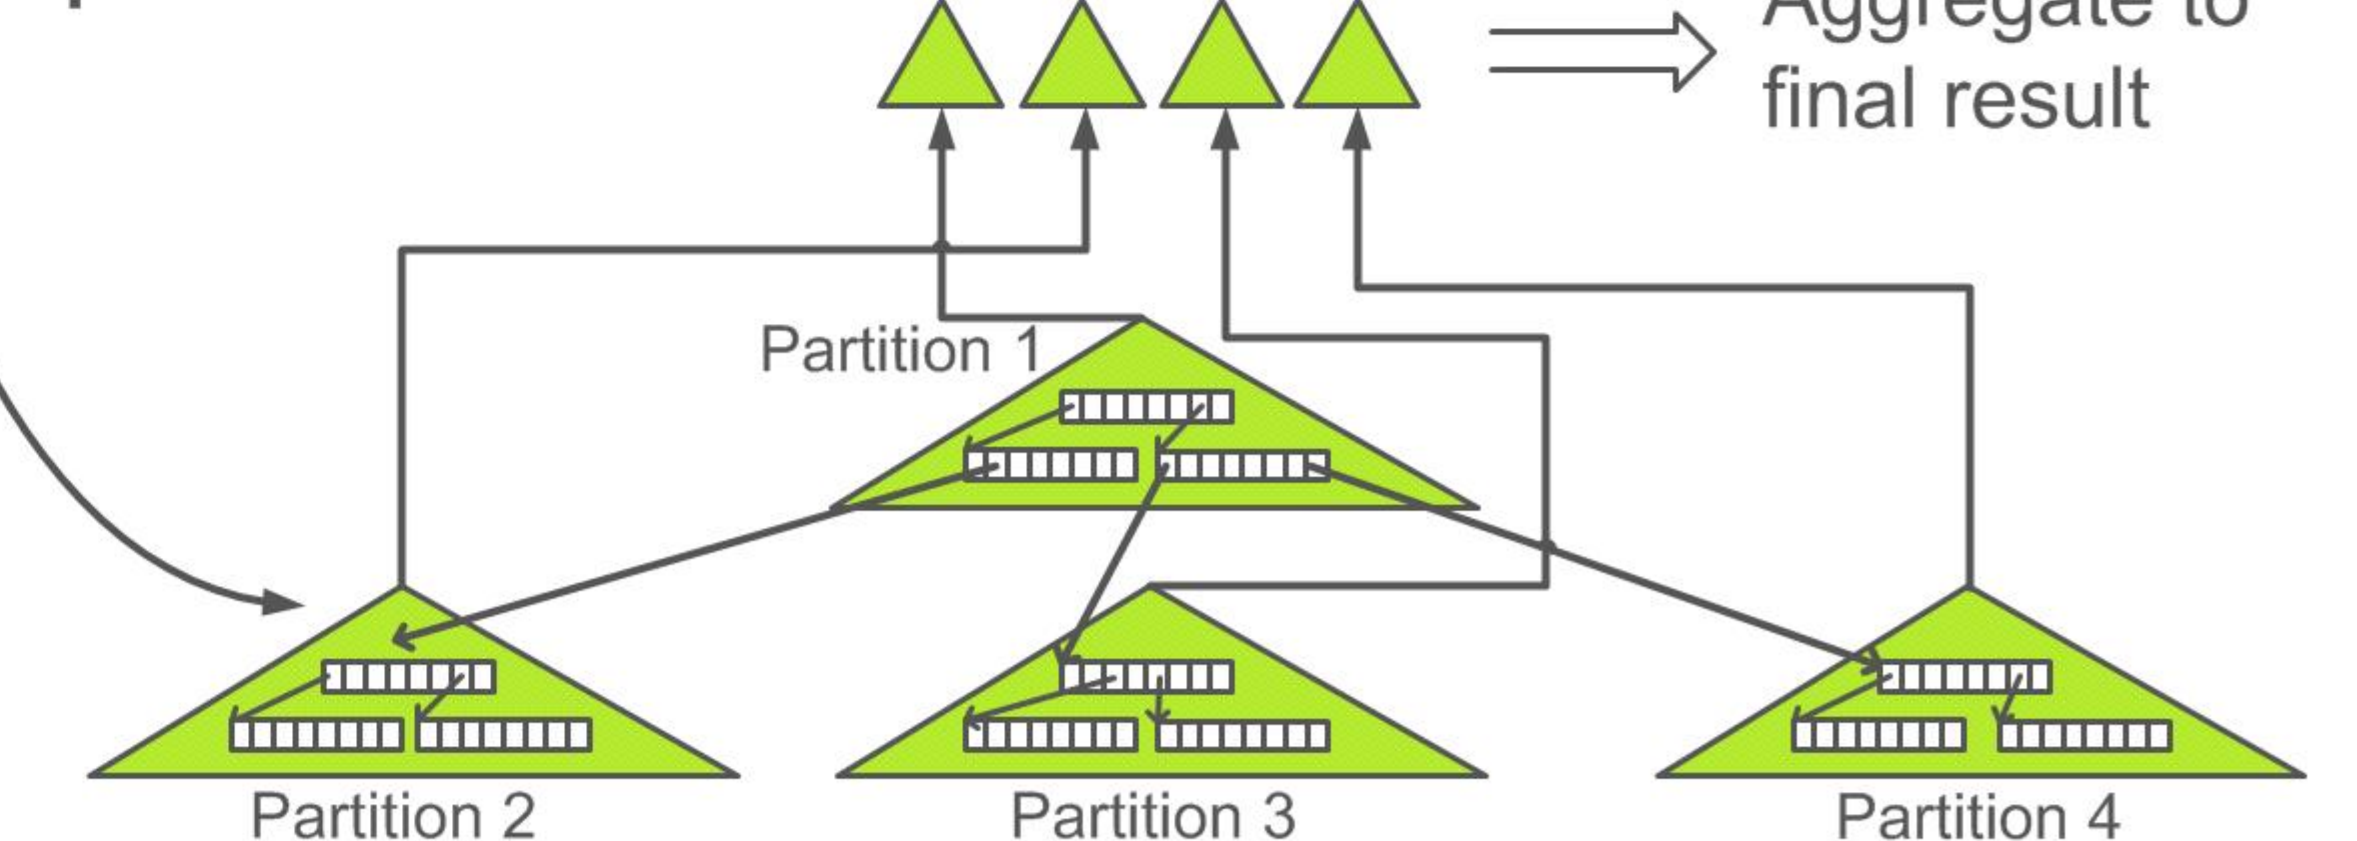


## Speculative Index Search\*

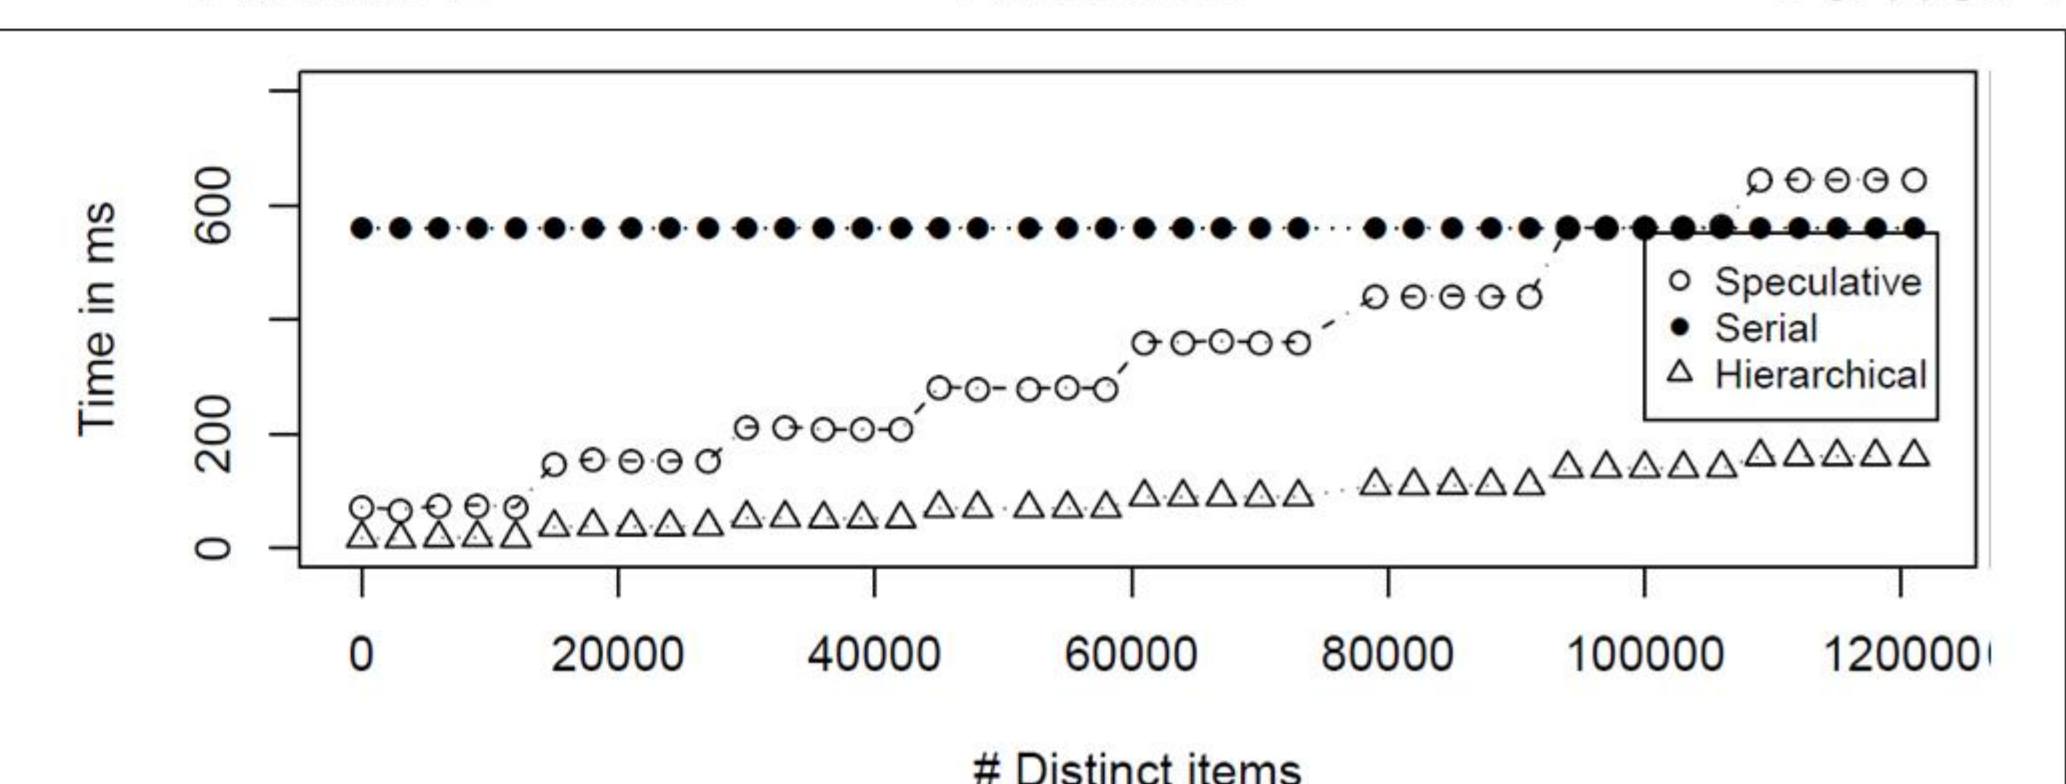
Serial search:



Speculative search:



### Evaluation



This work is supported by a Nvidia Professor Partnership grant

\*"GPU-Based Speculative Query Processing for Database Operations" - IN proceedings of First International Workshop on Accelerating Data Management Systems Using Modern Processor and Storage Architectures - In conjunction with VLDB2010

